

The Anatomy of a RAG Defense: Stress Tests, Failure Modes, and Design Lessons for Knowledge Poisoning

Aquib Raza Raja[†], *Student Member IEEE* and Muhammad Naveed Aman[†], *Senior Member IEEE*

[†]*School of Computing, University of Nebraska-Lincoln, Lincoln, NE, USA.*

Emails: araja2@huskers.unl.edu; naveed.aman@unl.edu

Abstract—[TODO: Write last, against the final result set.] The abstract must (i) lead with the stress-testing methodology contribution rather than with an attack result, (ii) name the four defense families studied and label the two proxies as proxies, (iii) state the evidence ceiling explicitly (oracle-level and development-set scale for the text-mutation results), and (iv) avoid any benchmark-wide rate.

Index Terms—retrieval-augmented generation, knowledge poisoning, defense evaluation, adversarial robustness, data integrity, trustworthy AI

I. INTRODUCTION

Motivation. [TODO: RAG adoption in production question answering and the resulting knowledge poisoning risk; corpus contents are now part of the trust boundary [1], [2].]

Gap. [TODO: Published defenses are evaluated almost exclusively against the attack distribution they were designed for, at the operating point their authors chose. A defense’s *stated* assumption and its *operative* assumption are not the same object, and only an assumption-targeted stress operator distinguishes them.]

Thesis. RAG poisoning defenses rest on different assumptions about how poisoned passages differ from benign ones: cross-passage embedding clustering, lexical and semantic density, learned feature separability, or answer aggregation. Stress tests constructed to target each assumption directly reveal distinct failure modes and distinct tradeoffs, and no single defense family is uniformly robust.

Contributions.

- 1) A defense-assumption stress-testing framework for poisoned RAG that separates oracle-space fragility from text-realizable evasion, from retrieval-preserving evasion, from answer-level outcome.
- 2) An oracle perturbation analysis exposing RAGDefender failure boundaries, with the mechanistic driver localized to Stage-2 top-pair collapse rather than to generic distribution matching. **[Completed: oracle result]**
- 3) Targeted text mutation families that stress distinct defense assumptions, with a measured asymmetry in which assumptions are text-reachable. **[Completed: development-set result]**

- 4) A two-regime evaluation of those mutations under fixed retrieved context and under full-retrieval replacement.

[Completed: development-set result]

- 5) RobustRAG-KW, a generation-time isolate-then-aggregate proxy, evaluated as an *orthogonal* defense layer rather than as a stronger one. **[Completed: RobustRAG-KW shortlisted scale-up]**

- 6) Design lessons for future RAG poisoning defenses and for how such defenses should be evaluated.

[Pending: larger GPT confirmation set] and **[Pending: toolkit ablation]** — the introduction must state, in the contribution list itself, which contributions are currently development-scale. Do not defer this to Section VIII.

Roadmap. [TODO: One short paragraph mapping Sections II–IX.]

II. BACKGROUND AND RELATED WORK

A. Retrieval-Augmented Generation

[TODO: Dense retrieval and RAG pipeline; the retrieved context as an untrusted input channel [1], [3], [4].]

B. Knowledge Poisoning and PoisonedRAG

[TODO: Corpus poisoning attacks; the PoisonedRAG threat model and its black-box LM-targeted variant, which is the attack this paper’s defenses are stressed against [2].]

C. Passage-Filtering Defenses

[TODO: Post-retrieval filtering: RAGDefender’s cross-passage similarity clustering [5]; FilterRAG’s per-passage frequency-density statistic and its learned ML-FilterRAG variant [6]. State the structural contrast: RAGDefender’s decision is *set-level* and therefore needs benign passages to exist in the retrieved set, while FilterRAG’s is *per-passage* and does not.]

D. Generation-Time and Aggregation-Based Defenses

[TODO: Isolate-then-aggregate generation defenses [7], which move the decision from passage filtering to answer aggregation.] Our RobustRAG-KW implementation is inspired by RobustRAG-style isolate-then-aggregate logic but is only a keyword/vote *proxy*; it is not a certified reproduction

of RobustRAG, and no certification result is claimed for it anywhere in this paper.

E. Robustness Evaluation and Stress Testing

[TODO: Semantics-preserving text attacks as the text-space operator family: TextFooler’s semantic-preservation and grammaticality constraints, BERT-Attack’s masked-LM fluency, and TextAttack’s sentence-encoder cosine constraints [8]–[10]. Distribution matching as a diagnostic tool: CORAL and MMD [11], [12].]

F. Positioning of This Work

This paper is a defense-assumption stress test, not simply a new attack paper. We do not claim a stronger attack than PoisonedRAG; we claim a way to find out what a defense actually depends on, and we apply it to four defense families. **[TODO: One sentence stating the gap precisely: no prior work separates oracle fragility from text-realizable evasion from retrieval-preserving evasion from answer-level outcome, on the same defenses and the same queries.]**

III. THREAT MODEL AND DEFENSES UNDER STUDY

A. Threat Model

Attacker capabilities. **[TODO: Black-box with respect to the generator; may author passage text and inject it into the corpus; knows the target query and the target incorrect answer; does *not* control retrieval, the defense implementation, or its thresholds.]**

Defender capabilities. **[TODO: Operates post-retrieval (RAGDefender, FilterRAG, ML-FilterRAG) or at generation time (RobustRAG-KW); sees only the retrieved passages and the query; has no poison labels at inference time.]**

Poison budget and retrieval setting. $N = 5$ poisoned passages per target query, retrieval depth $k = 10$ for the stress-testing regimes. **[TODO: Cross-reference the hyperparameter rationale in Sec. V-B.]**

Replacement-only mutation evaluation. When a mutation family is evaluated, the mutated passages *replace* the original poisoned passages. The poison budget is preserved at $N = 5$ per target query and no context ever contains both an original and its mutated counterpart. This matters because duplicating the poison would inflate every downstream number, and because retrieval must be rerun against a corpus holding one version of each poisoned passage, not two. **[Needs artifact verification: restate the one documented cross-query poison case, where a query retrieved poison authored for a different pool query, and point to the per-origin accounting introduced in Sec. IV-J]**

B. Defense Families

1) *RAGDefender*: **[TODO: Two-stage post-retrieval filter. Stage 1 estimates the number of adversarial passages from passage-level mean/median pairwise cosine similarity; Stage 2 ranks passages by an accumulated pair-frequency score over the top $(\hat{N}_{adv}/2)$ most similar pairs and**

removes the top $\hat{N}_{adv}$. Compress to at most six equations here; the full derivation stays in the external diagnostics, not in this manuscript.] Its operative dependence is that the poisoned passages form the densest high-similarity subgraph among the retrieved passages.

2) *FilterRAG*: **[TODO: Per-passage semantic frequency-density: the summed frequency of query and SLM-answer keywords matched into the passage, normalized by the passage’s unique word count; drop the passage when the score exceeds ϵ . Our implementation includes the paper-faithful semantic matching mode.] [Completed: paper-faithful semantic matching mode implemented]**

3) *ML-FilterRAG*: **[TODO: Learned passage classifier over FilterRAG-style density features plus fluency/likelihood features; remove when $P(\text{poison}) \geq t$.]** This is an *ML-FilterRAG-top-k proxy*: it scores the retrieved top- k set rather than reproducing the original top- s then filter then top- k pipeline. The label "proxy" must appear the first time the defense is named in every section.

4) *RobustRAG-KW*: **[TODO: Isolate-then-aggregate: answer each retrieved passage independently, then aggregate the isolated answers by keyword vote; abstain when the vote is contested.]** This is a RobustRAG-inspired keyword/vote *proxy*, not certified RobustRAG. It supports an orthogonality claim about *where* in the pipeline a defense acts, not a superiority claim over the passage filters.

C. Defense Assumptions

Table I states, for each family, the layer it acts at, the signal it reads, the assumption that signal encodes, and the stress operator this paper builds against that assumption. Fig. 2 gives the same relation as a conceptual map.

Caveats that must appear in this subsection, not only in Limitations. ML-FilterRAG is an ML-FilterRAG-top- k proxy. RobustRAG-KW is a RobustRAG-inspired keyword/vote proxy and not certified RobustRAG. Oracle perturbations act on embeddings or features in a frozen space and are not text-realizable by themselves: no perturbed vector has been shown to correspond to any real passage.

IV. STRESS-TESTING METHODOLOGY

A. Overview of the Experimental Pipeline

Fig. 1 shows the pipeline: poison authoring, retrieval, defense, generation, and metrics. Each stress operator enters at exactly one stage, and the stage it enters at determines what its result can license.

B. Evaluation Regimes

Table II defines the four regimes, what each holds fixed, what each therefore tests, and its current completion status.

Oracle embedding and feature stress tests. Poison representations are modified directly in the space the defense’s decision rule reads, with retrieval membership and passage text held fixed. Establishes which statistical assumption the rule operatively depends on; establishes no attack.

TABLE I
DEFENSE FAMILIES, THE SIGNAL EACH DECISION RULE READS, THE ASSUMPTION THAT SIGNAL ENCODES, AND THE STRESS PRESSURE APPLIED IN THIS PAPER.

Defense	Layer	Signal	Key assumption	Stress-test pressure
RAGDefender	Post-retrieval	Pairwise embedding similarity (set-level)	Poison forms the densest high-similarity subgraph	Clean-anchor and distribution-matching oracles; discourse-diverse rewrites
FilterRAG	Post-retrieval	Semantic frequency-density (per passage)	Poison is keyword-dense in query and answer terms relative to its length	Low-overlap targeted rewrites
ML-FilterRAG (proxy)	Post-retrieval	Learned density plus fluency features	Poison is separable in the learned feature space	Clean-reference rewrites; transfer from low-overlap rewrites
RobustRAG-KW (proxy)	Generation	Isolated per-passage answers, keyword vote	Poisoned passages do not command the aggregate vote	Poison-consensus contexts surviving retrieval

TABLE II
EVALUATION REGIMES, WHAT EACH CONTROLS, WHAT EACH TESTS, AND STATUS.

Regime	What is controlled	What it tests	Status
Oracle	Text and retrieval fixed; embeddings or features perturbed	Which statistical assumption the decision rule operates on	Completed (RAGDefender, ML-FilterRAG proxy)
Fixed context	Top- k membership and order fixed; poison text rewritten	Whether the assumption is reachable by a real passage	Completed on the development set (6 queries)
Full retrieval	Corpus substitution, replacement only, budget preserved	Whether rewritten poison survives rerun retrieval	Completed on 18 development-set bundles
Answer generation	Post-defense context passed to the generator	Whether the attacker answer is produced, or abstention occurs	Completed on 9 shortlisted cases

Fixed-context mutation. Poison passage *text* is rewritten and substituted into a retrieved context whose membership and order are held fixed by construction. Establishes that the assumption is reachable by a real passage; establishes nothing about retrieval.

Full-retrieval replacement. Mutated poison replaces original poison in the corpus and retrieval is rerun under a preserved poison budget. Establishes that the rewrite still retrieves.

Answer-generation ASR. Generation is run on the post-defense context and the target incorrect answer is scored. Establishes an answer-level outcome for the cases evaluated, and only for those cases.

C. Oracle Perturbations

All four operators act on a frozen embedding space with retrieval membership and passage text held fixed.

E1 clean-anchor interpolation. Each poisoned embedding is interpolated toward an assigned clean anchor, $\mathbf{z}' = \alpha\mathbf{z}_{\text{poison}} +$

$(1-\alpha)\mathbf{z}_{\text{anchor}}$, under several anchor-assignment strategies and a sweep over α . [TODO: Name the anchor strategies and the α grid.]

MMD. Kernel mean matching of the poisoned group toward the clean group, with a preservation penalty controlling how far embeddings may move.

CORAL-PCA and CORAL-ridge. Covariance alignment treating poisoned embeddings as source and clean embeddings as target, in a reduced-rank and a regularized form respectively.

These are not text-space attacks and not learned domain-adaptation models; they are formal embedding-space stress tests of a similarity-based decision rule.

D. GPT Targeted Mutation Families

Three hand-authored GPT rewrite families, one per targeted assumption. Each family rewrites all 5 poisoned passages of a query while preserving the attacker’s target answer.

RAGDefender-targeted discourse-diverse rewrites. Objective: reduce cross-passage similarity concentration so the poison set stops being the densest subgraph.

FilterRAG-targeted low-overlap rewrites. Objective: preserve the false claim while lowering query and answer keyword overlap and raising unique-word count.

ML-FilterRAG clean-reference rewrites. Objective: move the learned feature vector toward the clean feature distribution using clean passages as stylistic references.

[TODO: State that these rewrites are hand-authored via fixed GPT prompt packets on 6 development-set queries, that the prompts are hashed and committed, and that prompt design is therefore a development-set artifact — which is precisely why the confirmation set of Sec. IV-H is held out.]

E. Fixed-Context Evaluation

[TODO: Retrieved top- k membership and order reconstructed verbatim from the pilot artifacts; only the 5 poisoned passages’ text differs per family; clean passages byte-identical across all conditions. Report per-defense change in removed poison and in residual poison fraction.]

F. Full-Retrieval Replacement

[TODO: Mutated passages substituted into the corpus, retrieval rerun, replacement only, budget preserved. Report how many self-query poisoned passages return to the top- k , their rank range, and any cross-query poison retrieved.]

G. Answer Generation and RobustRAG-KW Evaluation

[TODO: Post-defense contexts passed to the generator; strict token-boundary matching for ASR; abstention detected by the same detector RobustRAG-KW uses, so that a refusal and a correct answer are never merged into one outcome. Cases are *shortlisted* for filter stress, so counts are reported, never rates.] [Do not overclaim: shortlist only]




Fig. 1. Stress-testing pipeline. A stress operator enters at exactly one stage, and the entry point determines what its result licenses. **[TODO: Redraw with the four regimes of Table II marked as gates between stages.]**

H. Confirmation-Set Protocol

[Pending: larger GPT confirmation set] The current 6-query targeted mutation work is treated as the development/prompt-design set. The planned confirmation set is 20 held-out HotpotQA queries at $k = 10$ and $N = 5$, with 3 mutation families $\times$ 2 bundles $\times$ 5 passages, i.e. 600 mutated passages. Query identifiers are fixed in advance and disjoint from the development set. **[TODO: Once the confirmation set exists, the development-set numbers in Sec. VI-C become prompt-design evidence and the confirmation-set numbers become the reported result. The prose in Sec. VI-C is written to make that substitution without restructuring.]**

I. Toolkit Mutation Ablation Protocol

[Pending: toolkit ablation] Planned as an ablation, after the confirmation set, asking whether generic library perturbations reproduce the failures produced by targeted GPT rewrites. The question is deliberately narrow: it is a control on the *specificity* of the targeted families, not a search for a stronger attack. **[TODO: Fix the toolkit, the constraint set, and the matched budget before running, so the comparison is a controlled substitution of the rewrite operator only.]**

J. Reproducibility Controls

- Cached retrieval and cached defense scoring, so a rerun reproduces prior numbers exactly rather than approximately.
- Hashed and committed GPT prompt packets for every mutation family.
- Replacement-only poison budget: $N = 5$ preserved in every mutated condition; no original-plus-mutated duplication.
- Held-out query identifiers fixed in advance for the confirmation set.
- No API calls in the test suite; every evaluation script consumes committed artifacts and loads models read-only.
- Self-query versus cross-query poison accounted separately, so a case that retrieves poison authored for a different pool query is never counted as if all retrieved poison targeted the case at hand.




Fig. 2. Conceptual map from defense families to the assumptions their decision rules encode. The edge from ML-FilterRAG to lexical/semantic density is the shared-feature path that predicts transfer of low-overlap rewrites. **[TODO: Annotate each edge with the stress operator that targets it.]**

V. EXPERIMENTAL SETUP

A. Dataset and Retrieval Setup

[TODO: HotpotQA (BEIR) with a dense retriever; state the encoder and scoring function, and that all stress-testing regimes use one dataset and one retriever [4], [13].]

B. Poisoning Budget and Top- k

We set $N = 5$ poisoned passages per target query to match the PoisonedRAG default setting [2]. We evaluate $k = 5$ and $k = 10$ to separate two retrieval regimes. The $k = 5$ setting corresponds to the PoisonedRAG default and represents a poison-saturated context in which the number of retrieved passages equals the number of injected poisoned passages. The $k = 10$ setting probes a mixed clean/poison regime where $k > N$, which PoisonedRAG’s own ablation identifies as the point where clean retrieved passages enter the context while poisoned passages are still usually retrieved. This mixed regime is particularly important for post-retrieval defenses such as RAGDefender and FilterRAG, because their filtering decisions depend on the composition of the retrieved top- k set. The $k = 5$ regime also violates RAGDefender’s stated requirement that at least one benign passage be retrieved, which is why $k = 10$ is the regime used for all stress-testing regimes above the nominal one.

C. Defense Implementations

[TODO: One short paragraph per family naming the concrete implementation and every proxy substitution: the small LM standing in for the paper’s generator inside FilterRAG, the perplexity model, the ML-FilterRAG-top- k pipeline substitution, and the keyword-vote form of RobustRAG-KW. These belong here, not only in Section VIII.]

D. Model and Generation Setup

[TODO: Generator model, decoding settings, and the fixed prompt template used for every answer-generation condition; note that RobustRAG-KW issues one isolated call per retrieved passage plus one aggregation step.]

E. Development Set versus Confirmation Set

The 6-query targeted mutation work is the *development set*: it is where the rewrite prompts were designed, so it cannot also serve as held-out confirmation. The planned 20-query dataset of Sec. IV-H is the *confirmation set*. **[Pending: larger GPT confirmation set]**

F. Thresholds and Operating Points

[TODO: FilterRAG ϵ and the fact that it was not retuned for the proxy small LM.] For ML-FilterRAG-top- k we use $t = 0.4$ throughout as the clean-preserving operating point, selected from the threshold sweep of Table VII and held fixed in every cross-defense comparison. **[Completed: development-set result]**

G. Metrics

- *Removed poison* and *removed clean*: counts of retrieved poisoned and clean passages the defense removes.
- *Residual poison fraction*: the poison share of the *surviving* context, not of the retrieved set. This definition differs from the intuitive reading and must be stated explicitly in the paper.
- *Strict ASR*: the attacker’s target answer appears in the generated answer under token-boundary matching.
- *Abstention*: the generator or aggregator declines to answer; counted separately from a correct answer.
- *Wrong-answer vote share* and *correct-answer vote share*: for RobustRAG-KW, the fraction of isolated per-passage answers voting for the attacker’s answer and for the gold answer.
- *Self-query versus cross-query poison*: retrieved poisoned passages authored for this query versus for a different pool query, reported separately.

H. Reporting Conventions

[TODO: State the known limitations of strict matching, in particular that token-boundary matching over-counts on yes/no answers and on answers that are prefixes of other tokens, and that a strict hit co-occurring with the correct answer is adjudicated rather than counted automatically. Report counts rather than rates wherever the cases were shortlisted.] **[Do not overclaim: shortlist only]**

VI. RESULTS

A. RQ1: What assumptions do existing defenses rely on?

[TODO: Read Table I and Fig. 2 as a result, not as background: the four families are separated not by strength but by the layer they act at and the signal they read, and that separation is what makes their failure modes disjoint.]

TABLE III
ORACLE PERTURBATIONS AGAINST RAGDEFENDER ON THE 6
ORIGINALLY SUCCESSFUL DEVELOPMENT-SET CASES. ORACLE-SPACE
EVIDENCE ONLY.

Perturbation	Failures	Main effect	Interpretation
E1 clean-anchor	6/6	Poison-poison count collapses	Breaks the operative assumption directly
MMD	Partial [Needs artifact verification: count]	Collapse only at weak preservation	Breaks it when the penalty permits large movement
CORAL-PCA	0	Distribution distance falls; top-pair structure intact	Distribution matching is not sufficient
CORAL-ridge	0	Largest displacement, no failures	Displacement magnitude is not the operative variable

B. RQ2: Can oracle perturbations expose RAGDefender failure boundaries?

Regime: oracle embedding space, text and retrieval membership held fixed. **[Completed: oracle result]**

E1 clean-anchor interpolation breaks all 6 of the 6 processable RAGDefender-success cases. MMD breaks them partially, and depends on the preservation penalty. CORAL-PCA and CORAL-ridge do not break them, despite reducing measured distribution distance and displacing poisoned embeddings further than the interventions that succeed. The mechanistic driver is therefore Stage-2 top-pair and pair-frequency collapse, not generic distribution matching: a defense can be moved a long way in distribution space without being moved at all in the direction its decision rule reads. **[TODO: State the query-strategy denominator explicitly next to the 6/6 figure, so a reader cannot mistake per-strategy units for per-query units.]** **[Needs artifact verification: the failure denominators reported across the oracle comparison artifacts are stated with different units (query-level versus query-strategy); recount from the batch CSV and fix one convention]** **[Needs artifact verification: these oracle results were produced against the legacy RAGDefender configuration; a paper-faithful configuration exists and has not been rerun through the full oracle sequence, so this subsection must either scope its claim to the evaluated configuration or wait for the rerun]**

This is oracle embedding-space evidence. It is not automatically a text-realizable attack, and Sec. VI-C shows that in this case it is not.

C. RQ3: Do targeted text rewrites weaken passage filters?

Regime: fixed retrieved context, poison text rewritten, development set of 6 queries. **[Completed: development-set result]**

Generic rewrites produce modest effects and serve as the negative control: rewriting per se does not evade these filters. Among the targeted families, the FilterRAG-targeted low-overlap family produces the strongest cross-defense effect,

TABLE IV

FIXED-CONTEXT CROSS-DEFENSE EFFECT OF THE TARGETED MUTATION FAMILIES ON THE DEVELOPMENT SET. VALUES ARE THE MEAN CHANGE IN REMOVED POISON PASSAGES (OF 5) RELATIVE TO THE UNMUTATED CONTEXT; NEGATIVE MEANS THE DEFENSE REMOVED FEWER POISONED PASSAGES. 6 QUERIES, ONE BUNDLE PER QUERY AND FAMILY.

Mutation family	RAGD.	FR	ML ($t=0.4$)	Interpretation
RAGDefender-targeted	0.00	-0.17	-0.67	Intended target unmoved
FilterRAG-targeted	-0.33	-1.17	-2.17	Strongest defense effect
ML-FilterRAG-targeted	+0.17	0.00	-0.17	Narrow, own-target only

generic-rewrite control from an earlier pilot produced a smaller mean effect on the ML-FilterRAG proxy than any targeted family. **[Needs artifact verification: restate the generic-control figure with its own query count, which differs from 6]**

weakening FilterRAG and transferring to the ML-FilterRAG-top- k proxy; the ML-FilterRAG clean-reference family moves only its own target; and the RAGDefender-targeted family does not move RAGDefender at all. The asymmetry is the finding: the density-based assumptions are text-reachable under hand-authored rewriting and RAGDefender’s clustering assumption, which the oracle of Sec. VI-B breaks reliably, is not.

[Pending: larger GPT confirmation set] These are development-set numbers on the 6 queries whose rewrite prompts were designed against these same defenses. They are evidence that the mechanism is reachable, not an estimate of how often it is reachable.

D. RQ4: Do rewritten poison passages survive retrieval?

Regime: full-retrieval replacement, corpus substitution, budget preserved. **[Completed: development-set result]**

Replacement-only full retrieval was completed on 18 mutation-bundle-1 cases. In every case the 5 self-query mutated poison passages returned to the top-10, so the rewrites of Sec. VI-C do not purchase filter evasion at the cost of retrievability. A minority of cases additionally retrieved poison authored for a different pool query; those passages are accounted separately and are never counted as self-query poison. **[Pending: larger GPT confirmation set]** This is development-scale evidence, not benchmark-wide evidence, and it does not establish that retrieval survival is near-total in general.

E. RQ5: Does generation-time aggregation change the outcome?

Regime: answer generation on shortlisted full-retrieval cases. **[Completed: RobustRAG-KW shortlisted scale-up]** **[Do not overclaim: shortlist only]**

Table VI reports the 9 shortlisted cases. RobustRAG-KW did not produce the attacker answer in 3 shortlisted cases where at least one passage-filtering defense did. Two of those three outcomes are abstentions and one is a clean defense, so the row counts of Def. and Abs. must be read separately: a refusal and a correct answer are not the same outcome for a user. RobustRAG-KW also failed on the cases where

TABLE V

FULL-RETRIEVAL REPLACEMENT ON THE DEVELOPMENT SET: 18 BUNDLES, REPLACEMENT ONLY, POISON BUDGET PRESERVED AT $N = 5$.

Mutation family	Cases	Self-query poison retrieved	Cross-query poison	Defense degradation
RAGDefender-targeted	6	5/5 all	[Needs artifact verification: count]	Reproduces fixed context
FilterRAG-targeted	6	5/5 all	[Needs artifact verification: count]	Strongest, as in fixed context
ML-FilterRAG-targeted	6	5/5 all	[Needs artifact verification: count]	Own target only

artifact verification: fill the per-family cross-query poison counts and the retrieved rank range from the full-retrieval artifacts; the aggregate retrieval-survival claim is verified, the per-family breakdown is not]

```
PLACEHOLDER fig:mutation_transfer_heatmap
mutation family (rows) x defense (columns),
cell = mean change in removed poison
source: evaluation_normalized/
normalized_cross_defense_failure_matrix.csv
```

Fig. 3. Cross-defense transfer of the targeted mutation families on the development set. Rows are mutation families, columns are defenses, and cells are the mean change in removed poison passages. **[TODO: Generate at publication resolution; consider folding into Table IV if the page budget tightens.]**

every isolated per-passages call independently voted for the poisoned answer, and it abstained on the majority of cases under contested votes.

The claim this supports is orthogonality, not superiority: aggregation reads a different signal than the passage filters, so it can hold where they fail and fail where they hold. The 9 cases were shortlisted for filter stress and are not a benchmark-wide sample; no rate may be computed from them.

F. RQ6: Do toolkit mutations reproduce the targeted GPT failures?

[Pending: toolkit ablation] **[TODO: One paragraph, written after the confirmation set, reporting whether generic library perturbations under a matched budget reproduce the per-defense effects of the targeted GPT families. A negative result here strengthens the specificity claim of Sec. VI-C and should be reported as such.]**

VII. DISCUSSION AND DESIGN LESSONS

A. Lesson 1: Clustering defenses depend on poison-poison structure

[TODO: Anchor: Sec. VI-B. The dependence is on the poison set being the densest subgraph, not on any

TABLE VI
ROBUSTRAG-KW GENERATION-TIME AGGREGATION COMPARED WITH PASSAGE-FILTERING DEFENSES ON SHORTLISTED FULL-RETRIEVAL MUTATION CASES.

Family	Query	Self/All poison	Passage filters	RRKW	s_{wrong}	s_{gold}	Reading
FR-tgt	Gibson/Zurrapapote gin	5/5	RD=ASR, FR=ASR, ML=ASR	ASR	0.90	0.10	Poison consensus
FR-tgt	Menges/Avakian occupation	5/5	RD=Def., FR=ASR, ML=ASR	ASR	0.50	0.00	Poison consensus
FR-tgt	Schmeichel IFFHS 1992	5/5	RD=Def., FR=Def., ML=Abs.	Abs.	0.50	0.20	Abstain (contested)
FR-tgt	Teide/Garajonay parks	5/5	RD=Def., FR=Def., ML=Def.	Abs.	0.30	0.40	Abstain (contested)
FR-tgt	Roth/Childers England	5/7	RD=Def., FR=Def., ML=ASR	Abs.	0.50	0.50	Abstain (contested)
FR-tgt	Ferocactus/Silene plants	5/6	RD=Def., FR=ASR, ML=ASR	Def.	0.20	0.80	Defends evasion
RD-tgt	Gibson/Zurrapapote gin	5/5	RD=Def., FR=ASR, ML=Def.	Abs.	0.90	0.10	Abstain (contested)
RD-tgt	Roth/Childers England	5/7	RD=Def., FR=Def., ML=Def.	Abs.	0.50	0.50	Abstain (contested)
ML-tgt	Roth/Childers England	5/7	RD=Def., FR=Def., ML=Def.	Abs.	0.50	0.50	Abstain (contested)

The 9 cases are a shortlisted, filter-stress sample, not a benchmark-wide rate. ASR means the target wrong answer was produced under strict matching; Def. means the target was not produced; Abs. denotes abstention. Self/All poison is self-query poison retrieved over all poison retrieved, so 5/7 marks a case that also retrieved cross-query poison.

```
PLACEHOLDER fig:retrieval_survival
per case: self-query mutated poison retrieved
in top- $k$ , with rank range;
cross-query poison marked separately
source: full_retrieval_pilot/
```

Fig. 4. Full-retrieval replacement survival on the development set: mutated self-query poison returning to the top-10 and its rank range, with cross-query poison marked separately. [TODO: Generate from the full-retrieval artifacts.]

```
PLACEHOLDER fig:robustrag_vote_distribution
per shortlisted case: wrong-answer vs
correct-answer vote share, with the
abstention band shaded
source: robustrag_kw_scaleup_
generation_results.csv
```

Fig. 5. RobustRAG-KW vote shares over the 9 shortlisted cases, with the contested region that produces abstention marked. Counts only; the shortlist does not support a rate. [Do not overclaim: shortlist only]

TABLE VII
ML-FILTERRAG-TOP- k PROXY THRESHOLD SWEEP ON A HELD-OUT HOTPOTQA 50-QUERY DRY RUN (10 HELD-OUT QUERY IDS, 20 (q, k) RECORDS, 118 RETRIEVED POISONED AND 32 RETRIEVED CLEAN PASSAGES), WITH THE FILTERRAG SEMANTIC-THRESHOLD REFERENCE ROW. $t = 0.4$ IS THE CLEAN-PRESERVING OPERATING POINT USED THROUGHOUT.

Method / threshold	Poison recall	Clean FPR	Residual poison
FilterRAG semantic, $\epsilon = 0.2$	0.983	0.781	0.222
ML-FilterRAG, $t = 0.35$	0.958	0.219	0.167
ML-FilterRAG, $t = 0.4$	0.949	0.156	0.182
ML-FilterRAG, $t = 0.50$	0.915	0.156	0.270
ML-FilterRAG, $t = 0.90$	0.678	0.000	0.543

global property of the poison distribution, which is why distribution matching leaves the defense intact.]

B. Lesson 2: Overlap-based filters are vulnerable to low-overlap support

[TODO: Anchor: Sec. VI-C. A false claim can be supported in text that is not keyword-dense in the query or the answer.]

C. Lesson 3: Learned filters inherit their feature assumptions

[TODO: Anchor: Sec. VI-C and Fig. 2. The ML-FilterRAG proxy inherits frequency-density as a feature, so transfer from low-overlap rewrites is a predicted consequence of the shared feature rather than a coincidence.]

D. Lesson 4: Aggregation is orthogonal but abstention-heavy

[TODO: Anchor: Sec. VI-E. Aggregation can resist some filter-evasion cases, fails under poison consensus, and frequently abstains under contested votes. State the safety-versus-coverage tradeoff plainly and do not present abstention as a defense success.]

E. Lesson 5: Oracle vulnerability does not imply a text-realizable attack

[TODO: Anchor: Sec. VI-B versus Sec. VI-C. The operator that most reliably breaks RAGDefender in embedding space is the one whose text-space approximation moved RAGDefender least. Acceptable framing: RAGDefender is fragile under oracle embedding perturbation and harder to break with natural text rewrites, and we report that gap rather than closing it.] [Pending: stronger RAGDefender text-space mutation operator]

F. Practical Implications for Future RAG Defenses

[TODO: Short paragraph: state assumptions explicitly and distinguish stated from operative; evaluate outside the attack distribution the defense was designed for; require retrieval preservation before claiming evasion and answer-level evidence before claiming an end-to-end outcome; report security and clean-context utility together, since a defense that empties the context has relocated the failure rather than fixed it; and combine layers that read different signals, because the failure modes documented here are disjoint.]

VIII. LIMITATIONS AND FUTURE WORK

Development-set size. The targeted mutation work covers 6 queries and 18 bundles, with one rewrite per poison slot per family, so per-family statistics are means over queries and not over repeated independent attempts.

HotpotQA-only current evidence. All stress-testing regimes use one dataset and one retriever; cross-dataset and cross-retriever generality is not measured.

Strict ASR limitations. Token-boundary matching overcounts, most visibly on yes/no answers, and a strict hit co-occurring with the correct answer requires adjudication.

RobustRAG-KW proxy limitation. A keyword/vote proxy, not certified RobustRAG; no certification result is claimed.

Answer-generation cost and shortlist bias. The 9 cases were shortlisted because passage filters were already under stress on them, so their outcome distribution is not that of the benchmark. **[Do not overclaim: shortlist only]**

Pending confirmation set. **[Pending: larger GPT confirmation set]** 20 held-out queries, 3 families, 2 bundles, 600 mutated passages, with the larger full-retrieval replacement run and the shortlisted answer generation on that set both pending.

Pending toolkit ablation. **[Pending: toolkit ablation]**

Additional datasets and defenses. **[TODO: Name the next dataset and the next defense family, with the reason each was chosen.]**

Stronger RAGDefender text-space mutations. **[Pending: stronger RAGDefender text-space mutation operator]** An operator with an explicit objective on the Stage-2 pair-frequency quantity, rather than on discourse diversity, is the natural next step; either outcome is informative, since success closes the gap of Lesson 5 and failure sharpens it.

Implementation fidelity. **[Needs artifact verification: state the proxy substitutions and any deviation of each defense implementation from its published description, including which RAGDefender configuration produced the oracle results of Sec. VI-B]**

IX. CONCLUSION

[TODO: Write last, in about four sentences. Restate the assumption-to-operator framework, the finding that top-pair collapse rather than distribution matching drives the clustering defense’s failure, the asymmetry between text-reachable and text-unreachable assumptions, and the orthogonality of generation-time aggregation. No new claims.]

ACKNOWLEDGMENT

[TODO: Funding and computing acknowledgment placeholder.]

REFERENCES

- [1] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, 2020.
- [2] W. Zou, R. Geng, B. Wang, and J. Jia, “PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models,” in *Proc. USENIX Security Symposium*, 2025.
- [3] V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih, “Dense passage retrieval for open-domain question answering,” in *Proc. Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [4] G. Izacard, M. Caron, L. Hosseini, S. Riedel, P. Bojanowski, A. Joulin, and E. Grave, “Unsupervised dense information retrieval with contrastive learning,” *Transactions on Machine Learning Research*, 2022.
- [5] S. Kim *et al.*, “Rescuing the unpoisoned: Efficient defense against knowledge corruption attacks on RAG systems,” in *Proc. Annual Computer Security Applications Conference (ACSAC)*, 2025, vERIFY full author list and page range.
- [6] K. Edemacu *et al.*, “Defending against knowledge poisoning attacks during retrieval-augmented generation,” in *arXiv preprint*, 2025, vERIFY venue and full author list.
- [7] C. Xiang, T. Wu, Z. Zhong, D. Wagner, D. Chen, and P. Mittal, “Certifiably robust RAG against retrieval corruption,” in *arXiv preprint*, 2024, vERIFY final venue.
- [8] D. Jin, Z. Jin, J. T. Zhou, and P. Szolovits, “Is BERT really robust? a strong baseline for natural language attack on text classification and entailment,” in *Proc. AAAI Conference on Artificial Intelligence*, 2020.
- [9] L. Li, R. Ma, Q. Guo, X. Xue, and X. Qiu, “BERT-ATTACK: Adversarial attack against BERT using BERT,” in *Proc. Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- [10] J. X. Morris, E. Lifland, J. Y. Yoo, J. Grigsby, D. Jin, and Y. Qi, “TextAttack: A framework for adversarial attacks, data augmentation, and adversarial training in NLP,” in *Proc. EMNLP: System Demonstrations*, 2020.
- [11] B. Sun, J. Feng, and K. Saenko, “Return of frustratingly easy domain adaptation,” in *Proc. AAAI Conference on Artificial Intelligence*, 2016.
- [12] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf, and A. Smola, “A kernel two-sample test,” *Journal of Machine Learning Research*, vol. 13, 2012.
- [13] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning, “HotpotQA: A dataset for diverse, explainable multi-hop question answering,” in *Proc. Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.